

In [1]:

```
"""# I. Preparing the dataset """  
#1 Importing essential libraries  
import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
import statsmodels.api as sm  
  
#2 Importing the dataset  
dataset = pd.read_csv('Salary_Data.csv')
```

In [2]:

```
#3 Classify dependent and independent variables
X = dataset.iloc[:, :-1].values #independent variable YearsofExperience
y = dataset.iloc[:, -1].values #dependent variable salary

print("\nIndependent Variable (Experience):\n", X)
print("\nDependent Variable (Salary):\n", y)

plt.scatter(X, y, alpha=1)
```

Independent Variable (Experience):


```
[[ 1.1]
 [ 1.3]
 [ 1.5]
 [ 2. ]
 [ 2.2]
 [ 2.9]
 [ 3. ]
 [ 3.2]
 [ 3.2]
 [ 3.7]
 [ 3.9]
 [ 4. ]
 [ 4. ]
 [ 4.1]
 [ 4.5]
 [ 4.9]
 [ 5.1]
 [ 5.3]
 [ 5.9]
 [ 6. ]
 [ 6.8]
 [ 7.1]
 [ 7.9]
 [ 8.2]
 [ 8.7]
 [ 9. ]
 [ 9.5]
 [ 9.6]
 [10.3]
 [10.5]]
```

Dependent Variable (Salary):

```
[ 39343.  46205.  37731.  43525.  39891.  56642.  60150.  544
 45.  64445.
 57189.  63218.  55794.  56957.  57081.  61111.  67938.  6602
 9.  83088.
 81363.  93940.  91738.  98273. 101302. 113812. 109431. 10558
 2. 116969.
112635. 122391. 121872.]
```

Out[2]:

```
<matplotlib.collections.PathCollection at 0x1f5bd99f688>
```



In [3]:

```
#4 Creating training set and testing set
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X ,y, test_size = 1/3,random_state = 0)

print("\n\nTraining Set :\n-----\n")
print("X = \n", X_train)
print("y = \n", y_train)

print("\n\nTest Set :\n-----\n")
print("X = \n",X_test)
print("y = \n", y_test)
```

Training Set :

-----

```
X =
[[ 2.9]
 [ 5.1]
 [ 3.2]
 [ 4.5]
 [ 8.2]
 [ 6.8]
 [ 1.3]
 [10.5]
 [ 3. ]
 [ 2.2]
 [ 5.9]
 [ 6. ]
 [ 3.7]
 [ 3.2]
 [ 9. ]
 [ 2. ]
 [ 1.1]
 [ 7.1]
 [ 4.9]
 [ 4. ]]

y =
[ 56642.  66029.  64445.  61111. 113812.  91738.  46205. 1218
 72.  60150.
 39891.  81363.  93940.  57189.  54445. 105582.  43525.  3934
 3.  98273.
 67938.  56957.]
```

Test Set :

-----

```
X =
[[ 1.5]
 [10.3]
 [ 4.1]
 [ 3.9]
 [ 9.5]
 [ 8.7]
 [ 9.6]
 [ 4. ]
 [ 5.3]
 [ 7.9]]

y =
[ 37731. 122391.  57081.  63218. 116969. 109431. 112635.  557
```

```
94. 83088.  
101302.]
```

In [4]:


```
#5 Train the Regressor with training set  
model = sm.OLS(y_train, X_train).fit()  
  
#6 predict the outcome of test sets  
y_Pred = model.predict(X_test)
```

In [5]:

```
#7 Mapping the Regression Line  
plt.scatter(X, y, alpha=1)  
plt.plot(X_test, y_Pred, c = 'blue')
```

Out[5]:

[<matplotlib.lines.Line2D at 0x1f5be23c3c8>]



In [6]:

```
#8 Calculating the Accuracy of the predictions  
from sklearn import metrics  
print("Prediction Accuracy = ", metrics.r2_score(y_test, y_Pred))  
  
# Print out the statistics  
print(model.summary())
```



Prediction Accuracy = 0.8111465166115929

### OLS Regression Results

```

=====
Dep. Variable:          y      R-squared (uncentered):
                                0.965
Model:                  OLS    Adj. R-squared (uncentered):
                                0.963
Method:                 Least Squares    F-statistic:
525.3
Date:                   Sat, 14 Sep 2019    Prob (F-statistic):
2.64e-15
Time:                   19:21:29    Log-Likelihood:
-219.36
No. Observations:      20    AIC:
440.7
Df Residuals:          19    BIC:
441.7
Df Model:              1
Covariance Type:       nonrobust
=====

```

```

=====
              coef      std err          t      P>|t|
[0.025      0.975]
-----
x1          1.375e+04    600.077     22.920     0.000     1.25e+04
=====

```

```

=====
Omnibus:              0.809    Durbin-Watson:
2.456
Prob(Omnibus):        0.667    Jarque-Bera (JB):
0.454
Skew:                 -0.361    Prob(JB):
0.797
Kurtosis:             2.843    Cond. No.
1.00
=====

```

### Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.



In [7]:

```
#9 Some more insights
print ('MAE:', metrics.mean_absolute_error(y_Pred, y_test))
print ('RMSE:', np.sqrt(metrics.mean_squared_error(y_Pred, y_test)))
print ('R-Squared:', metrics.r2_score(y_Pred, y_test))
```

MAE: 10828.326955554774  
 RMSE: 12581.660968849912  
 R-Squared: 0.9016135955873352

In [8]:

```
#10 Comparing Actual and Predicted Salaries for the test set
print("\nActual vs Predicted Salaries \n-----\n")
print("Actual :\n ", y_test)
print("Predicted :\n ", y_Pred)
```

Actual vs Predicted Salaries  
 -----

Actual :  
 [ 37731. 122391. 57081. 63218. 116969. 109431. 112635. 55  
 794. 83088.  
 101302.]  
 Predicted :  
 [ 20630.47442402 141662.59104493 56389.96342565 53639.2335  
 0245  
 130659.67135212 119656.75165931 132035.03631372 55014.598464  
 05  
 72894.34296487 108653.8319665 ]